

The Local Housing-Market Effects of Act 22/60 Investor Purchases: Data and Empirical Design

August 2026

1 Research question

When a tax-incentivized migrant (an Act 22/60 “Individual Investor” decree holder) buys a residential property in Puerto Rico, what happens to the housing market immediately around that property — to prices, and to who transacts?

2 Data construction

2.1 Identifying investor properties

- **Decree holders.** Government lists of individual grantees: 3,181 (Act 22) and 2,672 (Act 60) names; 5,770 distinct first–last pairs pooled.
- **CRIM cadastral registry.** Name searches against the CRIM parcel layer (via its ArcGIS REST service) matched 26,673 unique parcels, each with full attributes, centroid coordinates, and the parcel’s *most recent* sale (price, date, buyer, seller). CRIM stores only the latest transaction per parcel.
- **Karibe property registry.** Registry name-search results; cadastral numbers extracted from registrarial descriptions and resolved in CRIM provide a second, largely non-overlapping source.
- **Match confidence.** Name→parcel matching is noisy for common names. A name returning *exactly one* property is treated as high-confidence (“unique match”): 810 parcels via CRIM, 755 via Karibe (surname-validated, geocoded), union **1,306 parcels** (259 cross-validated in both). The event studies use this base.
- **Coverage.** High-confidence identification covers 22.5% of searched names (26.7% including a validated multi-match recovery tier), versus 49–61% declared homeownership in the Act 22 annual reports. The gap is dominated by LLC/trust purchases (invisible to name search), name-format failures, and registry availability; the identified sample skews toward single-property, personal-name buyers.

2.2 Events

An **event** is a dated investor purchase: unique-match parcels with a valid CRIM sale date, collapsing same-day, same-location purchases (condo units bought together) into one event, giving 1,256 events (1,200 in the decree era, ≥ 2012 , used in estimation). Purchases cluster heavily: the median event has 25 other events within 1 km (158 are fully isolated), so contamination is measured per event and used in robustness.

2.3 Outcomes: nearby sales

For each event, every CRIM parcel within 1,000 m with a recorded sale: 3.47M sale $\times$ event pairs over 328K distinct parcels, each carrying distance, sale price/date/parties, property characteristics (lot size, assessed structure value, condo-unit and vacant-land indicators), and 2020 census tract. *Rings*: 0–250 m (treated), 250–400 m (buffer, dropped), 400–1,000 m (control). *Filters*: nominal transfers ($< \$10k$), implausible dates, and sales of investor-matched parcels are excluded. *Caveat*: CRIM keeps only each parcel’s last sale, so earlier sales are censored by later resales; the within-window, near-versus-far design mitigates this, and Figure 5 tests it.

2.4 Placebo events

1,500 purchases by non-corporate *individual* buyers in the investor price band (\$285K–\$3.9M, the p25–p95 of investor purchases), located within 250 m of a real event site (where the sales pool has full coverage) but more than three years apart in time. Far rings are truncated at 750 m for both the placebo run and a matched real-event baseline.

2.5 Buyer/seller origin proxy

Party names are classified against the Census 2010 surnames file (`pcthispanic`), taking the maximum across name tokens (robust to CRIM’s inconsistent name ordering; appropriate under the Puerto Rican two-surname convention). Names with `pcthispanic` ≤ 20 are “non-Hispanic-named” (mainland proxy), ≥ 70 “Hispanic-named” (local proxy); corporate, ambiguous, and unmatched names are left unclassified, and the raw score is retained for re-thresholding. *Validity*: 67% of investor purchases have non-Hispanic-named buyers versus 10% of surrounding market sales. *Caveat*: the proxy captures *name* origin — stateside Puerto Ricans and other Hispanic-named mainlanders classify as local, so mainland penetration and displacement effects are understated.

2.6 Summary statistics

Table 1 summarizes the estimation samples: the ring-design sale sample, the Red Atlas tract price panel, the HMDA purchase panel, and the 2010 PRCS baseline tract characteristics used for balance and robustness.

Table 1: **Summary Statistics** This table presents summary statistics for the paper’s main samples. Panel A: market sales within 1km of an identified decree-era investor purchase (ring design estimation sample: gap ring, junk/nominal prices, and investor-parcel sales excluded; events 2012+); name-classification indicators are defined over sales whose parties’ surnames classify. Panel B: Red Atlas tract×year panel (transaction-count-weighted annual mean prices); identified purchases reported over the 235 treated tracts. Panel C: HMDA tract×year panel of originated first-lien owner-occupied 1–4 family purchases, 2012–2024. Panel D: 2010 PRCS tract characteristics (2020 boundaries, population-weighted).

	Mean (1)	Std. Dev. (2)	P25 (3)	Median (4)	P75 (5)	Observations (6)
Panel A: Ring-design sales ($\leq 1\text{km}$ of an event)						
Sale price (\$000s)	1,872.3	61,984.0	116.6	240.0	475.0	971,720
Log sale price	12.35	1.27	11.67	12.39	13.07	971,720
1{non-Hispanic-named buyer}	0.169	0.375	0.000	0.000	0.000	660,570
1{Hisp. seller $\rightarrow$ non-Hisp. buyer}	0.115	0.320	0.000	0.000	0.000	500,309
1{non-Hisp. seller $\rightarrow$ Hisp. buyer}	0.097	0.295	0.000	0.000	0.000	500,309
Panel B: Tract price panel (Red Atlas)						
Tract annual mean price (\$000s)	148.5	122.9	91.2	122.3	168.5	16,344
Tract annual transactions	25.5	24.0	11.0	19.0	33.0	16,344
Identified purchases (treated tracts)	5.1	13.7	1.0	1.0	3.0	235
Panel C: HMDA purchases (2012–2024)						
Hispanic-borrower purchases	11.6	12.3	4.0	8.0	15.0	11,350
Non-Hispanic-borrower purchases	0.195	0.819	0.000	0.000	0.000	11,350
Mean borrower income (\$000s)	54.7	79.8	34.8	43.9	58.0	11,100
Panel D: 2010 PRCS tract characteristics						
Median household income (\$)	20,242	10,759	13,704	17,102	23,700	875
Median home value (\$)	119,534	54,345	88,700	104,150	134,675	874
Median gross rent (\$)	461	181	352	436	541	863
Poverty rate	0.457	0.165	0.347	0.473	0.576	877
BA+ share (25+)	0.208	0.128	0.120	0.175	0.257	878
Renter share	0.291	0.155	0.187	0.251	0.358	877
Vacancy rate	0.162	0.088	0.105	0.147	0.199	877
Seasonal-home share	0.032	0.060	0.006	0.017	0.036	877
Born-in-mainland share	0.050	0.025	0.031	0.047	0.064	878
Population	4,020	2,008	2,634	3,990	5,320	936

3 Empirical design

3.1 Panel construction

Because CRIM records one (the last) sale per parcel, sales are aggregated into an event×ring×period panel. Let e index events, $r \in \{\text{near, far}\}$ rings, and t calendar periods (years in the robustness

suite). The cell outcome is the mean log price among the N_{ert} transactions in cell (e, r) in period t :

$$\bar{y}_{ert} = \frac{1}{N_{ert}} \sum_{j \in (e, r, t)} \ln P_j, \quad (1)$$

where P_j is the sale price of transaction j . Treatment is absorbing and turns on for the near cell in the event period t_e :

$$D_{ert} = \mathbf{1}\{r = \text{near}\} \times \mathbf{1}\{t \geq t_e\}, \quad (2)$$

so far cells are never treated and serve as clean controls. In the tract-fixed-effects specification, $\ln P_j$ is first residualized on census tract dummies at the sale level ($\ln P_j - \hat{\gamma}_{\tau(j)}$) and the residual is aggregated instead; results are unchanged.

3.2 Local projections difference-in-differences (LP-DiD)

Following Dube, Girardi, Jordà and Taylor (2023), the dynamic effect at each horizon h is estimated from a separate long-difference regression on the stacked cell panel. For post-treatment horizons $h = 0, 1, \dots, H$:

$$\bar{y}_{er, t+h} - \bar{y}_{er, t-1} = \beta_h^{\text{LP-DiD}} \Delta D_{ert} + \delta_t^h + \varepsilon_{ert}^h, \quad (3)$$

and for pre-treatment horizons $h = 2, \dots, Q$ (with $h = 1$ the normalization):

$$\bar{y}_{er, t-h} - \bar{y}_{er, t-1} = \beta_{-h}^{\text{LP-DiD}} \Delta D_{ert} + \delta_t^h + \varepsilon_{ert}^h, \quad (4)$$

where $\Delta D_{ert} = D_{ert} - D_{er, t-1}$ indicates that the near cell of event e enters treatment in period t , and δ_t^h are horizon-specific calendar-period effects.

Each horizon- h regression is estimated on the restricted “clean” sample

$$\left\{ (e, r, t) : \underbrace{\Delta D_{ert} = 1}_{\text{newly treated near cells}} \quad \text{or} \quad \underbrace{D_{er, s} = 0 \ \forall s}_{\text{never-treated far cells}} \right\}, \quad (5)$$

which is the **nevertreated** option: the counterfactual change $\bar{y}_{er, t+h} - \bar{y}_{er, t-1}$ for a newly treated near cell is identified from far cells over the same calendar window, and no already-treated observation is ever used as a control (the negative-weighting problem of two-way fixed effects under staggered timing cannot arise). $\beta_h^{\text{LP-DiD}}$ is a variance-weighted average of event-specific effects; standard errors are clustered by cell. The identifying assumption is within-micro-neighborhood parallel trends: absent the purchase, prices of homes transacting 0–250 m from the site would have evolved like those 400–1,000 m away, conditional on calendar-time effects — shocks common to the micro-area (Hurricane María, the 2020–21 in-migration wave) difference out by construction, and equations (4) provide the testable pre-trend restrictions $\beta_{-h} = 0$.

Two scope conditions follow from the ring structure. First, the near–far contrast differences out any effect *common to the full 1-km neighborhood*: the design identifies the hyper-local increment around a purchase site, not the neighborhood-level effect of aggregate investor inflows, which

requires the complementary tract-level staggered design using the tract $\times$ month price panel (Section 5). Second, because purchases cluster at the enclave scale, other investor purchases in dense areas fall disproportionately inside the near ring itself; the estimated contrast there measures the cumulative local dose of investor activity rather than the effect of a single purchase (see Figure 5).

3.3 Robustness to the estimation approach

The cell-mean panel is a construction choice; the ring literature (Linden and Rockoff 2008; Campbell, Giglio and Pathak 2011) instead estimates on individual transactions: stacked sale-level regressions of log price on near $\times$ post (or near $\times$ event-time dummies) with event $\times$ ring and event $\times$ calendar-year fixed effects, standard errors clustered by event. The two implementations also weight differently — transactions versus events — so agreement is informative rather than mechanical. Re-estimating the baseline this way yields a pooled near $\times$ post effect of +7.09 log points (s.e. 1.58), and +6.72 (1.32) with hedonic controls (log lot size, sub-unit and vacant-land indicators), against the cell-level LP-DiD baseline of +6.95 (1.81); the event-study paths coincide horizon by horizon (Figure 1).




Figure 1: **Estimation-approach robustness.** The near-ring event study estimated three ways: the cell-level LP-DiD baseline; sale-level stacked OLS with event $\times$ ring and event $\times$ calendar-year fixed effects (the ring literature’s convention), clustered by event; and the sale-level specification with hedonic controls. Rings 0–250 m vs 400–1,000 m. The three series coincide within sampling error at every horizon, and the agreement of transaction-weighted and event-weighted estimates indicates the pooled effect is not driven by a few high-volume events.

In the annual robustness suite $Q = 4$ and $H = 3$; the baseline is also estimated at quarterly ($Q = H = 8$) and half-yearly resolutions and with the Callaway–Sant’Anna (2021) group-time estimator $ATT(g, t) = \mathbb{E}[\bar{y}_t - \bar{y}_{g-1} \mid G_g = 1] - \mathbb{E}[\bar{y}_t - \bar{y}_{g-1} \mid \text{never treated}]$, aggregated to event

time; results are similar across estimators and resolutions, and identical with and without tract fixed effects.

4 The eight figures

All figures plot LP-DiD event-study coefficients $\{\beta_{-Q}, \dots, \beta_H\}$ from equations (3)–(4) at annual resolution (pre-window 4, post-window 3), with 95% confidence intervals, cells clustered. Period -1 is the omitted base; the dashed vertical line marks the treatment boundary.




Figure 2: **Baseline event study.** Near-ring (0–250 m) versus far-ring (400–1,000 m) effects on log sale price by year relative to the investor purchase. The headline price bump; flat pre-period coefficients are the design’s core credibility check, and the post-period path shows how quickly the premium appears and whether it persists.



Figure 3: **Investor vs. placebo purchases.** The real event study overlaid with placebo events (non-investor individual buyers in the investor price band, same micro-areas, >3 years from the real event; far rings matched at 400–750 m). A flat placebo says nearby prices respond to *the investor*, not to any high-value transaction; because unmatched investors contaminate the placebo toward finding an effect, a null placebo is conservative evidence of investor specificity.





Figure 4: **Spatial gradient.** Near-ring (0–250 m) and gap-ring (250–400 m) series, each against the same far control. A genuine hyper-local spillover should decay with distance (near > gap > 0); near $\approx$ gap would instead indicate a neighborhood-wide shock partially shared by the control ring.



Each series: that group's near rings vs never-treated controls; annual lpdid, 95% CIs clustered by cell.
Groups by total identified purchases within 1km of the event (the event itself included), matching the tract design's 1-4 / 5+ split.

Figure 5: **Estimated contrast by local investor concentration.** Separate event studies for events with 1–4 vs. 5+ total identified purchases within 1 km (the event itself included), matching the tract design's dose split. Because purchases cluster at the enclave scale (condo buildings, resort communities), other events in dense areas fall disproportionately *inside* the near ring, so the near–far contrast there reflects cumulative local investor exposure rather than a single purchase. Low-concentration events show no price effect (post–pre -2.7 , s.e. 3.5); 5+ areas show $+6.4$ (s.e. 1.5): single-purchase channels (comps entry, one renovation, signaling) appear weak on their own, and the price response is a phenomenon of investor *concentration* — the same 1–4/5+ pattern as the tract design. Pre-period coefficients are flat in both groups, inconsistent with investor waves selecting into already-appreciating pockets.




Figure 6: **Censoring robustness.** Baseline overlaid with events from 2018 onward, for which little time exists for post-period sales to be overwritten by later resales. Agreement with the full sample indicates that CRIM's last-sale-only censoring is not manufacturing the bump.




Figure 7: **Composition checks.** The identical design with property characteristics as outcomes: log assessed structure value (top left), log lot size (top right), condo-unit share (bottom left), vacant-land share (bottom right). Flat lines imply the price bump is appreciation on a stable transaction mix; movements would reveal the stock-upgrading margin of gentrification, changing the interpretation of Figure 2.




Figure 8: **Who transacts: buyers vs. sellers.** Effects on the share of nearby sales with non-Hispanic-named buyers and, separately, sellers (Census 2010 surname classification; corporate and ambiguous names excluded). Buyer share rising faster than seller share means the area is net-importing mainland owners after an investor arrives — the composition shift price indices cannot see.




Figure 9: **Transition types.** Effects on the shares of transactions that are Hispanic-seller→non-Hispanic-buyer versus non-Hispanic-seller→non-Hispanic-buyer (among sales where both parties classify). Rising islander→mainlander transitions indicate conversion of locally held housing to incomers (the displacement margin); rising mainlander→mainlander indicates churn within an already-converted segment — together distinguishing “investors lead conversion” from “investors cluster where conversion already occurred.”




Figure 10: **Net compositional inflow.** The two transition series combined into one displacement measure: $1\{\text{Hispanic seller, non-Hispanic buyer}\}$ minus $1\{\text{non-Hispanic seller, Hispanic buyer}\}$ per classified sale. A sale changes the ownership composition of the stock only when the parties' classes differ, so within-class churn (including mainlander→mainlander resales) nets out; a positive post-event effect is net conversion of the nearby stock from Hispanic-named to non-Hispanic-named owners. Pre-period flat (-0.5pp , s.e. 1.0); the net conversion rate rises to $+2.5\text{pp}$ and $+3.8\text{pp}$ per sale in years 2–3; post–pre $+2.4\text{pp}$ (s.e. 0.9, $p < 0.01$) per classified sale.

5 The tract-level design and reconciliation (Design 2)

5.1 Data and specification

The tract-level design asks what the ring design structurally cannot: whether investor purchases move *whole neighborhoods*, including the component common to treated areas that near–far contrasts difference out. The outcome panel is the Red Atlas tract×month transaction file (918 of 981 2020-vintage tracts, 2008–2025), aggregated to tract×year with the annual price defined as the transaction-count-weighted mean of monthly tract means (composition-noise reduction; the simple mean is a robustness variant). Treatment timing comes from the dated decree-era events: a tract is treated from the year of its first identified investor purchase (235 treated tracts; 1,200 events), with never-treated tracts as controls. Estimation is annual LP-DiD (pre-window 4, post-window 3) on $100 \times \ln(\text{price})$, with and without county fixed effects.

Table 2 compares treated and never-treated tracts on 2010 PRCS characteristics, all pre-determined relative to Act 22. Treated tracts are richer (median income +\$5,700), higher-value (+\$39,000), more educated, less poor, and markedly more seasonal-home- and mainlander-oriented — exactly the profile of investor selection into coastal-amenity areas. Residualizing out municipio fixed effects shrinks the gaps only modestly (15–30%; all remain significant): treated tracts are distinct *within* municipios, not just across them. Selection on 2010 observables is a fact about *where* investors go; the specification ladders below (Table 3) absorb municipio-year shocks directly, and the dose exhibits show the resulting effect heterogeneity explicitly.

Table 2: **2010 Characteristics by Tract Treatment Status** This table presents summary statistics for 2010 PRCS (ACS 2006–2010 5-year) tract characteristics — all pre-determined relative to Act 22 (January 2012) — separated by tracts that ever receive an identified decree-era investor purchase (treated) and tracts that never do. Column (7) reports the raw mean difference between treated and never-treated tracts and column (8) the difference after residualizing out municipio (county) fixed effects, with heteroskedasticity-robust p -values in parentheses. Dollar variables in 2010 dollars; shares in $[0,1]$. 2010 tracts are mapped to 2020 boundaries by largest land overlap, population-weighted. Statistical significance levels are indicated as follows: *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

	Treated Tracts			Never-Treated Tracts			Mean Diff.	Adj. Mean Diff.
	Mean	Std. Dev.	Observations	Mean	Std. Dev.	Observations	[(1) – (4)]	[(1) – (4)]
	(1)	(2)	(3)	(4)	(5)	(6)	p -value	p -value
Median household income (\$)	24,522	13,955	218	18,821	9,039	657	5,701*** (0.000)	4,774*** (0.000)
Median home value (\$)	148,938	81,270	218	109,762	36,943	656	39,176*** (0.000)	28,126*** (0.000)
Median gross rent (\$)	515	206	216	442	169	647	73*** (0.000)	62*** (0.000)
Poverty rate	0.400	0.169	218	0.476	0.160	659	-0.076*** (0.000)	-0.065*** (0.000)
BA+ share (25+)	0.270	0.168	218	0.188	0.104	660	0.082*** (0.000)	0.071*** (0.000)
Renter share	0.286	0.155	218	0.292	0.156	659	-0.006 (0.594)	-0.035*** (0.003)
Vacancy rate	0.196	0.119	218	0.152	0.072	659	0.044*** (0.000)	0.037*** (0.000)
Seasonal-home share	0.058	0.102	218	0.024	0.032	659	0.034*** (0.000)	0.030*** (0.000)
Born-in-mainland share	0.057	0.032	218	0.047	0.022	660	0.010*** (0.000)	0.007*** (0.002)
Population	4,549	1,935	218	3,859	2,004	718	690*** (0.000)	972*** (0.000)

5.2 Results

Prices in treated tracts rise by a pooled +7.6 log points (s.e. 1.8) after the first investor purchase, with flat pre-period coefficients; adding county fixed effects leaves the estimate unchanged (+7.5), and the simple-mean construction gives +6.1 (Figure 11). Splitting treated tracts by investor concentration reproduces the dose pattern at tract scale: tracts with five or more identified purchases (44 tracts) show +18.3 (2.9), tracts with one to four show +3.8 (2.1) (Figure 12). Transaction *volume* falls in treated tracts, but its pre-period coefficients are significantly negative, so parallel trends fails for that outcome and we treat it as descriptive only.

Table 3 subjects the tract-level estimates — prices, borrower income, and purchase counts by ethnicity — to a common specification ladder, all estimated with the paper’s baseline pre-mean-differenced LP-DiD: (1) calendar-year effects and (2) county $\times$ year effects, so identification is within municipio-year. The price effect moves from 6.4 to 5.7 (both $p < 0.01$), borrower income from 8.1

to 7.8 ($p < 0.01$), non-Hispanic purchase entry ($100 \times \text{asinh}$) from 12.7 to 12.2 ($p < 0.01$), and Hispanic purchase counts stay a precise null. Table 4 runs the ladder on the ring design’s cell panel (specification (1) reproduces the baseline estimates exactly), adding in column (3) house-characteristic controls from the transaction data — cell means of log assessed structure value, log lot size, sub-unit share, and vacant-land share, pre-mean differenced like the outcome. The price effect goes 5.3, 4.2, 3.0 (all $p < 0.01$; the R-squared doubles to 0.53), so a constant-composition repricing of roughly +3 log points survives holding fixed what transacts; net ownership conversion is untouched (2.4, 2.5, 2.3pp, all $p < 0.01$). Municipio-year shocks account for little of any estimate. We do not add 2010 PRCS covariate interactions to these ladders: ACS suppression drops $\approx 9\%$ of cells concentrated in exactly the low-population resort tracts where effects are largest, so that specification conflates sample composition with adjustment (in a fixed complete-controls sample the ten interactions move the ring price estimate by only ≈ 0.5 log points). Table 5 collects the concentration results of both designs on the common 1–4/5+ dose convention, and Table 6 reports the placebo contrast in table form.

Table 3: **Tract-level treatment effects are robust to county-year shocks** This table reports Treated coefficients from the paper’s baseline estimator — the pre-mean-differenced LP-DiD (Dube et al. 2023; $H=3$, $k=4$, clean-control sample: newly treated tracts at onset plus never-treated tracts) — under two specifications per outcome: (1) calendar-year effects; (2) county $\times$ year effects, so identification is within municipio-year. Tract fixed effects are subsumed by the long-differencing. Panel A: tract log price (Red Atlas, count-weighted annual mean) and log mean purchase-borrower income (HMDA, 2012–2024), both scaled by 100. Panel B: purchase origination counts by borrower ethnicity as $100 \times \text{asinh}$ (Poisson does not compose with pre-mean differencing; the Poisson event studies appear in the figures). Robust standard errors are clustered at the tract level and reported in parentheses. Statistical significance levels are indicated as follows: *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

Panel A: prices and borrower income				
	<i>100 $\times$ Log(Price)</i>		<i>100 $\times$ Log(Borrower Income)</i>	
	(1)	(2)	(1)	(2)
Treated	6.44*** (1.41)	5.66*** (1.59)	8.10*** (1.99)	7.81*** (2.23)
Observations	11,714	11,681	8,319	8,299
R-squared	0.153	0.257	0.054	0.160
Tract differencing (LP-DiD)	Yes	Yes	Yes	Yes
Calendar-year effects	Yes	No	Yes	No
County-year effects	No	Yes	No	Yes
Panel B: purchase originations by borrower ethnicity (100$\times$asinh)				
	<i>Hispanic Purchases</i>		<i>Non-Hispanic Purchases</i>	
	(1)	(2)	(1)	(2)
Treated	4.04 (5.36)	2.97 (5.75)	12.75*** (3.00)	12.16*** (3.09)
Observations	9,414	9,392	9,414	9,392
R-squared	0.088	0.167	0.019	0.111
Tract differencing (LP-DiD)	Yes	Yes	Yes	Yes
Calendar-year effects	Yes	No	Yes	No
County-year effects	No	Yes	No	Yes

Table 4: **Ring-design treatment effects: county-year shocks and house-characteristic controls**

This table subjects the ring design’s estimates to a specification ladder, using the paper’s baseline estimator throughout: the pre-mean-differenced LP-DiD on the event×ring cell panel ($H=3$, $k=4$; clean-control sample: newly treated near cells at onset plus never-treated cells; specification (1) reproduces the baseline table exactly). Specifications: (1) calendar-year effects; (2) county×year effects (the county is the event’s municipio), so identification is within municipio-year; (3) adds house-characteristic controls from the transaction data — cell-year means of log assessed structure value, log lot size (cabida), sub-unit share, and vacant-land share, each pre-mean differenced exactly like the outcome — so the column (3) Treated coefficient is a constant-composition price effect. Cell fixed effects are subsumed by the long-differencing. Outcomes: cell mean log sale price (×100) and the net Hispanic-to-non-Hispanic ownership conversion rate per classified sale (percentage points). Robust standard errors are clustered at the cell level and reported in parentheses. Statistical significance levels are indicated as follows: *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

	$100 \times \text{Log}(\text{Price})$			<i>Net H→NH Conversion (pp)</i>		
	(1)	(2)	(3)	(1)	(2)	(3)
Treated	5.32*** (1.40)	4.23*** (1.34)	3.02*** (1.13)	2.37*** (0.88)	2.47*** (0.86)	2.28*** (0.85)
Observations	12,379	12,122	12,058	12,035	11,772	11,723
R-squared	0.154	0.343	0.531	0.047	0.168	0.174
Cell differencing (LP-DiD)	Yes	Yes	Yes	Yes	Yes	Yes
Calendar-year effects	Yes	No	No	Yes	No	No
County-year effects	No	Yes	Yes	No	Yes	Yes
House-characteristic controls	No	No	Yes	No	No	Yes

Table 5: **Price Effects by Investor Concentration: Ring and Tract Designs** This table reports LP-DiD price effects by investor concentration under both designs, on the tract design’s dose convention. Columns (1)-(2): ring design, events split by total identified purchases within 1km of the event (the event itself included); each group’s near rings (0–250m) vs never-treated control rings, cell-level annual LP-DiD. Columns (3)-(4): tract design, treated tracts split by identified purchases in the tract, each group vs never-treated tracts. The Treated coefficient is the post–pre effect from the single pre-mean-differenced LP-DiD regression (Dube et al. 2023), whose SE, R-squared, and Observations are reported. All outcomes $100 \times \log$ price; SEs clustered by cell (ring) or tract. Statistical significance levels are indicated as follows: *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

	<i>Ring design</i>		<i>Tract design</i>	
	1–4 purch. (1)	5+ purch. (2)	1–4 purch. (3)	5+ purch. (4)
Treated	-2.74 (3.45)	6.41*** (1.50)	4.13*** (1.59)	16.30*** (2.54)
Observations	2,854	9,525	11,670	11,476
R-squared	0.074	0.217	0.153	0.152
Unit differencing (LP-DiD)	Yes	Yes	Yes	Yes
Calendar-year effects	Yes	Yes	Yes	Yes
Never-treated controls	Yes	Yes	Yes	Yes

Table 6: **Investor Purchases vs Placebo Luxury Purchases** Placebo events are 1,500 purchases by individual non-investor buyers in the investor price band, drawn from the same micro-areas and more than three years from any real event; both series use far rings matched at 400–750m. Columns (1)–(2): cell mean log sale price ($\times 100$). Columns (3)–(4): net Hispanic-to-non-Hispanic ownership conversion per classified sale (percentage points; $1\{\text{Hisp. seller, non-Hisp. buyer}\} \text{ minus } 1\{\text{non-Hisp. seller, Hisp. buyer}\}$). Responses unique to investor purchases rule out generic luxury-transaction effects on both prices and neighborhood ownership composition. The Treated coefficient is the post–pre effect from a single LP-DiD regression via pre-mean differencing (Dube, Girardi, Jordà and Taylor 2023): dependent variable $\frac{1}{4} \sum_{h=0}^3 y_{t+h} - \frac{1}{4} \sum_{\tau=t-4}^{t-1} y_{\tau}$ on treatment entry with calendar-time effects, clean-control sample; its SE, R^2 and Observations come from that regression. Coefficients are $100 \times \log$ points (shares in pp). SEs clustered by unit (cell/tract). Statistical significance levels are indicated as follows: *** $p < 0.01$, ** $p < 0.05$, * $p < 0.10$.

	$100 \times \text{Log}(\text{Price})$		$\text{Net } H \rightarrow NH \text{ Conversion (pp)}$	
	Investor (1)	Placebo (2)	Investor (3)	Placebo (4)
Treated	4.43*** (1.46)	-0.33 (1.40)	2.17** (0.90)	-0.28 (0.87)
Observations	12,036	7,932	11,425	7,680
R-squared	0.109	0.258	0.023	0.025
Calendar-year effects	Yes	Yes	Yes	Yes
Event $\times$ ring cell differencing	Yes	Yes	Yes	Yes
Hedonic controls	No	No	No	No

5.3 Reconciliation: the whole vs the sum of the parts

Because the ring design delivers a distance gradient, it implies a *prediction* for each treated tract: the stock-value-weighted average of $\beta(d)$ over the tract’s parcels — the tract-level effect if micro-spillovers were the whole story. Comparing observed tract-DiD estimates to these predictions decomposes the aggregate (Figure 13): tracts with 1–4 purchases sit at or below their prediction (observed +3.8 vs predicted +5.4) — micro-spillovers fully account for them — while the 44 concentration tracts exceed their prediction by roughly 11 log points (observed +18.3 vs predicted +7.5). A market-level channel beyond proximity operates only where investors concentrate.

5.4 Aggregation

Integrating the estimated gradient over the exposure distribution of the island’s improved housing stock (60.0% of parcels, 63.1% of assessed value, lie within 2.5 km of an identified event) yields an island-wide average price effect of +2.1 to +3.0%, or \$6.3–8.0B of housing value at band-specific market/assessed ratios — roughly 6–8% of Puerto Rico’s post-2019 boom (FHFA purchase-only index: +64% 2019–2024). Adding the concentration excess in the 44 high-dose tracts ($\approx \$4.3\text{B}$) raises the identified program footprint to $\approx \$11\text{--}12\text{B}$, about 10–11% of the boom. At municipio




Figure 11: **Tract-level event study: log house price.** Annual LP-DiD on the Red Atlas panel, 235 treated tracts vs never-treated controls; base and county-FE series are indistinguishable and pre-trends are flat.




Figure 12: **Tract-level event study by investor concentration.** Tracts with 5+ identified purchases versus tracts with 1–4, each against never-treated controls: the concentration result at tract scale.




Figure 13: **Observed tract effects vs ring-implied predictions.** Observed: LP-DiD pooled post estimates by dose group. Predictions: stock-value-weighted distance-gradient effects over each treated tract's parcels. The all-treated and 1–4-purchase groups sit on their predictions; the 5+ group's excess (≈ 11 log points) is the concentration channel.

level, using municipio repeat-sales indices as the local boom denominators, footprints in the top investor municipios are much larger (Table 7).

Table 7: **Program Footprint in the Most-Exposed Municipios** This table reports each municipio’s post-2019 housing boom and the share of it attributable to identified investor purchases under three accountings. All columns share one valuation convention: each parcel at assessed value times its tract’s median market/assessed ratio (2019+ sales; municipio and island fallbacks). The dollar boom is the municipio’s market stock times $1 - 1/(1 + g)$, the portion of current value attributable to repeat-sales growth g since 2019; Uplift is the synthesis in dollars. Shares of boom connect exactly: *Ring*, *Total* integrates the distance gradient over all parcels and *5+ tracts* restricts to parcels in 5+-purchase tracts; *Tract DiD* applies the dose-specific tract effects to treated-tract stock (*Total*: all treated tracts; *5+ tracts*: the high-dose subset). The synthesis *5+ differential* is Tract DiD (5+) minus Ring (5+) — the tract-wide repricing beyond what proximity predicts — and *Synthesis*, *Total* = Ring (Total) + the 5+ differential, so the ring layer is never double-counted.

Municipio	Local boom '19–'24			Synthesis		Ring design		Tract DiD	
	Growth (1)	\$ (2)	Uplift (3)	Total (4)	5+ diff. (5)	Total (6)	5+ tracts (7)	Total (8)	5+ tracts (9)
Rincón	+148%	\$1.1B	\$0.34B	31.1%	17.7%	13.4%	12.7%	30.7%	30.4%
Dorado	+177%	\$3.7B	\$0.97B	25.9%	14.4%	11.5%	10.3%	25.0%	24.8%
San Juan	+89%	\$17.8B	\$4.29B	24.1%	9.0%	15.1%	6.0%	17.8%	15.1%
Río Grande	+115%	\$2.0B	\$0.48B	24.0%	15.7%	8.4%	6.1%	23.5%	21.8%
Carolina	+91%	\$7.6B	\$1.32B	17.2%	6.8%	10.4%	4.7%	12.2%	11.6%
Guaynabo	+124%	\$5.4B	\$0.84B	15.5%	2.3%	13.2%	1.7%	8.6%	4.0%
Humacao	+121%	\$2.5B	\$0.39B	15.5%	7.2%	8.3%	5.1%	13.7%	12.4%

All aggregates are lower bounds: exposure counts only the 1,200 name-identified purchases, excluding LLC and unmatched investors; the capitalization step applies transaction-price effects to the full stock; and event effects predating 2019 enter the numerator while the boom clock starts in 2019.

5.5 Mortgage originations: volumes are a convincing null

HMDA loan-level data for Puerto Rico (2018–2024; 105,559 originations, aggregated to 2020-vintage tract×year) show no detectable response of mortgage lending in treated tracts: Poisson (PPML) estimates with tract and year fixed effects are statistically zero for owner-occupied purchases (−0.5%, s.e. 5.4), non-owner-occupied purchases (+3.0%, 11.3), rate/term and cash-out refinancings, and total originations, and are unchanged with county×year fixed effects. Consistent with decree investors transacting predominantly in cash, the program’s effects appear in prices and buyer composition rather than local credit expansion.

5.6 Who can still buy: borrower composition in treated tracts

Splitting purchase originations by HMDA’s self-reported borrower ethnicity (Hispanic or Latino vs. not; joint and unavailable excluded) sharpens the displacement interpretation. Since decree investors buy in cash, HMDA is effectively a census of the *financed* — overwhelmingly local, 94.6% Hispanic — side of the market. To get adequate pre-periods we extend the panel back to 2012 using

CFPB’s historic HMDA files, on a population that is consistent across the 2017/2018 reporting change: originated first-lien, owner-occupied, 1–4 family home-purchase loans (2012–2024; annual volumes are smooth across the seam). Three results (Figures 14–15). First, borrower *income* rises after treatment — pooled +9.6% (s.e. 1.9), reaching +10.0% at $h \geq 3$ — and with up to six pre-years of data (displayed with the endpoint binned at -4) the pre-period is now flat (-3.7 , -1.1 , -2.4 , individually and jointly insignificant, with no run-up into treatment; the apparent pre-drift in the short 2018–2024 panel was an artifact of its truncated window). Second, the income shift is concentrated where investors concentrate: +17.6% (s.e. 4.9) in 5+-purchase tracts vs. +7.9% (s.e. 2.0) in 1–4-purchase tracts. Third, on the extensive margin, Hispanic-borrower purchase counts are unchanged (pooled +0.6%, s.e. 4.5) while *non-Hispanic* borrower purchases rise +37% (s.e. 14) from a thin base — financed mainlander entry is now detectable, not just cash entry. Together: local buyers as a group keep transacting, but the marginal financed buyer moves up the income distribution, most strongly in high-dose tracts — the affordability-displacement margin, measured on self-reported ethnicity and income rather than the surname proxy.




Figure 14: **Income of purchase borrowers, 2012–2024 panel.** Effect on $100 \times \ln(\text{mean borrower income})$ among first-lien owner-occupied purchase originations in treated tracts; all classified borrowers and Hispanic borrowers only. No pre-trend (endpoint binned at -4 pools horizons back to -6); pooled +9.6% (s.e. 1.9).

Financed mainlander entry

The non-Hispanic borrower series make the entry margin visible on its own terms (Figures 16–17; thin base — a few hundred non-Hispanic-borrower loans per year island-wide — hence the wide intervals). Purchase originations by non-Hispanic borrowers rise +37% (s.e. 14) pooled and +41% (s.e. 19) at $h \geq 3$ in treated tracts, with noisy but non-trending pre-periods: decree-era in-movers are




Figure 15: **Borrower income by investor concentration, 2012–2024 panel.** Each series: that treated group vs. never-treated tracts. The income shift is twice as large in 5+-purchase tracts (+17.6%, s.e. 4.9) as in 1-4-purchase tracts (+7.9%, s.e. 2.0) — the same concentration gradient as prices.

not exclusively cash buyers — a financed segment follows the same geography.¹ Their refinancings rise even more steeply — +70% (s.e. 17) pooled, +176% (s.e. 34) at $h \geq 3$ — consistent with newcomers refinancing recently acquired properties as the areas appreciate. (The corresponding *Hispanic* refinancing series is not usable as evidence: its pre-period climbs steadily into treatment with no post-treatment kink, so the incumbent equity-access margin is unidentified and we do not report it.)

6 Standing caveats

1. **Last-sale censoring** (CRIM): pre-period sales are survivors that never resold; addressed by design and Figure 5.
2. **Site selection**: investors choose appreciating micro-areas; estimates are local spillovers *conditional on* location choice, not “the effect of the Act.”

¹This magnitude reconciles with the ring design’s buyer-composition effect once the two bases are compared. The CRIM non-Hispanic-*named* buyer share rises +2.5pp (s.e. 0.7, PMD post–pre) on a pre-treatment near-ring mean of 18.9% — nominally +13%. But the surname proxy inflates the base: self-reported non-Hispanic borrowers are only 1.6% of financed purchases in treated tracts pre-treatment, and non-Hispanic-*named* buyers are 10% of the general market against 67% among confirmed investors, so much of the 18.9pp proxy base is Hispanic-heritage locals with non-Hispanic-classifying surnames — a component inert to treatment. If the treatment-responsive (true mainlander) share of the proxy base is roughly a third of it (6–7pp), the +2.5pp effect is +35–40% of the responsive base, matching the HMDA estimate almost exactly. Both sources thus indicate mainlander purchasing rose by roughly one-third of its pre-treatment level — CRIM measuring all transactions at ring geography, HMDA the financed tail at tract geography.




Figure 16: **Purchase originations by non-Hispanic borrowers, 2012–2024 panel.** Poisson event study, treated vs. never-treated tracts. Pooled +37% (s.e. 14).




Figure 17: **Refinancings by non-Hispanic borrowers, 2012–2024 panel.** Same population (first-lien owner-occupied 1–4 family; all refi types — the cash-out split does not exist before 2018). Pooled +70% (s.e. 17), rising to +176% (s.e. 34) at $h \geq 3$.

3. **Match selection:** identified investors skew toward single-property, personal-name buyers; entity (LLC/trust) purchases are missing.
4. **Name proxy:** understates mainland presence; corporate and ambiguous names are excluded from shares.
5. **Assessed values** are on CRIM's 1957 basis — valid as cross-sectional controls and composition outcomes, never as market values.