

The Local Housing-Market Effects of Act 22/60 Investor Purchases: Data and Empirical Design

August 2026

1 Research question

When a tax-incentivized migrant (an Act 22/60 “Individual Investor” decree holder) buys a residential property in Puerto Rico, what happens to the housing market immediately around that property — to prices, and to who transacts?

2 Data construction

2.1 Identifying investor properties

- **Decree holders.** Government lists of individual grantees: 3,181 (Act 22) and 2,672 (Act 60) names; 5,770 distinct first–last pairs pooled.
- **CRIM cadastral registry.** Name searches against the CRIM parcel layer (via its ArcGIS REST service) matched 26,673 unique parcels, each with full attributes, centroid coordinates, and the parcel’s *most recent* sale (price, date, buyer, seller). CRIM stores only the latest transaction per parcel.
- **Karibe property registry.** Registry name-search results; cadastral numbers extracted from registral descriptions and resolved in CRIM provide a second, largely non-overlapping source.
- **Match confidence.** Name→parcel matching is noisy for common names. A name returning *exactly one* property is treated as high-confidence (“unique match”): 810 parcels via CRIM, 755 via Karibe (surname-validated, geocoded), union **1,306 parcels** (259 cross-validated in both). The event studies use this base.
- **Coverage.** High-confidence identification covers 22.5% of searched names (26.7% including a validated multi-match recovery tier), versus 49–61% declared homeownership in the Act 22 annual reports. The gap is dominated by LLC/trust purchases (invisible to name search), name-format failures, and registry availability; the identified sample skews toward single-property, personal-name buyers.

2.2 Events

An **event** is a dated investor purchase: unique-match parcels with a valid CRIM sale date, collapsing same-day, same-location purchases (condo units bought together) into one event, giving 1,256 events (1,200 in the decree era, ≥ 2012 , used in estimation). Purchases cluster heavily: the median event has 25 other events within 1 km (158 are fully isolated), so contamination is measured per event and used in robustness.

2.3 Outcomes: nearby sales

For each event, every CRIM parcel within 1,000 m with a recorded sale: 3.47M sale $\times$ event pairs over 328K distinct parcels, each carrying distance, sale price/date/parties, property characteristics (lot size, assessed structure value, condo-unit and vacant-land indicators), and 2020 census tract. *Rings*: 0–250 m (treated), 250–400 m (buffer, dropped), 400–1,000 m (control). *Filters*: nominal transfers ($< \$10k$), implausible dates, and sales of investor-matched parcels are excluded. *Caveat*: CRIM keeps only each parcel’s last sale, so earlier sales are censored by later resales; the within-window, near-versus-far design mitigates this, and Figure 5 tests it.

2.4 Placebo events

1,500 purchases by non-corporate *individual* buyers in the investor price band (\$285K–\$3.9M, the p25–p95 of investor purchases), located within 250 m of a real event site (where the sales pool has full coverage) but more than three years apart in time. Far rings are truncated at 750 m for both the placebo run and a matched real-event baseline.

2.5 Buyer/seller origin proxy

Party names are classified against the Census 2010 surnames file (`pcthispanic`), taking the maximum across name tokens (robust to CRIM’s inconsistent name ordering; appropriate under the Puerto Rican two-surname convention). Names with `pcthispanic` ≤ 20 are “non-Hispanic-named” (mainland proxy), ≥ 70 “Hispanic-named” (local proxy); corporate, ambiguous, and unmatched names are left unclassified, and the raw score is retained for re-thresholding. *Validity*: 67% of investor purchases have non-Hispanic-named buyers versus 10% of surrounding market sales. *Caveat*: the proxy captures *name* origin — stateside Puerto Ricans and other Hispanic-named mainlanders classify as local, so mainland penetration and displacement effects are understated.

3 Empirical design

3.1 Panel construction

Because CRIM records one (the last) sale per parcel, sales are aggregated into an event $\times$ ring $\times$ period panel. Let e index events, $r \in \{\text{near}, \text{far}\}$ rings, and t calendar periods (years in the robustness

suite). The cell outcome is the mean log price among the N_{ert} transactions in cell (e, r) in period t :

$$\bar{y}_{ert} = \frac{1}{N_{ert}} \sum_{j \in (e, r, t)} \ln P_j, \quad (1)$$

where P_j is the sale price of transaction j . Treatment is absorbing and turns on for the near cell in the event period t_e :

$$D_{ert} = \mathbf{1}\{r = \text{near}\} \times \mathbf{1}\{t \geq t_e\}, \quad (2)$$

so far cells are never treated and serve as clean controls. In the tract-fixed-effects specification, $\ln P_j$ is first residualized on census tract dummies at the sale level ($\ln P_j - \hat{\gamma}_{\tau(j)}$) and the residual is aggregated instead; results are unchanged.

3.2 Local projections difference-in-differences (LP-DiD)

Following Dube, Girardi, Jordà and Taylor (2023), the dynamic effect at each horizon h is estimated from a separate long-difference regression on the stacked cell panel. For post-treatment horizons $h = 0, 1, \dots, H$:

$$\bar{y}_{er, t+h} - \bar{y}_{er, t-1} = \beta_h^{\text{LP-DiD}} \Delta D_{ert} + \delta_t^h + \varepsilon_{ert}^h, \quad (3)$$

and for pre-treatment horizons $h = 2, \dots, Q$ (with $h = 1$ the normalization):

$$\bar{y}_{er, t-h} - \bar{y}_{er, t-1} = \beta_{-h}^{\text{LP-DiD}} \Delta D_{ert} + \delta_t^h + \varepsilon_{ert}^h, \quad (4)$$

where $\Delta D_{ert} = D_{ert} - D_{er, t-1}$ indicates that the near cell of event e enters treatment in period t , and δ_t^h are horizon-specific calendar-period effects.

Each horizon- h regression is estimated on the restricted “clean” sample

$$\left\{ (e, r, t) : \underbrace{\Delta D_{ert} = 1}_{\text{newly treated near cells}} \quad \text{or} \quad \underbrace{D_{er, s} = 0 \ \forall s}_{\text{never-treated far cells}} \right\}, \quad (5)$$

which is the **nevertreated** option: the counterfactual change $\bar{y}_{er, t+h} - \bar{y}_{er, t-1}$ for a newly treated near cell is identified from far cells over the same calendar window, and no already-treated observation is ever used as a control (the negative-weighting problem of two-way fixed effects under staggered timing cannot arise). $\beta_h^{\text{LP-DiD}}$ is a variance-weighted average of event-specific effects; standard errors are clustered by cell. The identifying assumption is within-micro-neighborhood parallel trends: absent the purchase, prices of homes transacting 0–250 m from the site would have evolved like those 400–1,000 m away, conditional on calendar-time effects — shocks common to the micro-area (Hurricane María, the 2020–21 in-migration wave) difference out by construction, and equations (4) provide the testable pre-trend restrictions $\beta_{-h} = 0$.

Two scope conditions follow from the ring structure. First, the near–far contrast differences out any effect *common to the full 1-km neighborhood*: the design identifies the hyper-local increment around a purchase site, not the neighborhood-level effect of aggregate investor inflows, which

requires the complementary tract-level staggered design using the tract $\times$ month price panel (Section 5). Second, because purchases cluster at the enclave scale, other investor purchases in dense areas fall disproportionately inside the near ring itself; the estimated contrast there measures the cumulative local dose of investor activity rather than the effect of a single purchase (see Figure 5).

3.3 Robustness to the estimation approach

The cell-mean panel is a construction choice; the ring literature (Linden and Rockoff 2008; Campbell, Giglio and Pathak 2011) instead estimates on individual transactions: stacked sale-level regressions of log price on near $\times$ post (or near $\times$ event-time dummies) with event $\times$ ring and event $\times$ calendar-year fixed effects, standard errors clustered by event. The two implementations also weight differently — transactions versus events — so agreement is informative rather than mechanical. Re-estimating the baseline this way yields a pooled near $\times$ post effect of +7.09 log points (s.e. 1.58), and +6.72 (1.32) with hedonic controls (log lot size, sub-unit and vacant-land indicators), against the cell-level LP-DiD baseline of +6.95 (1.81); the event-study paths coincide horizon by horizon (Figure 1).




Figure 1: **Estimation-approach robustness.** The near-ring event study estimated three ways: the cell-level LP-DiD baseline; sale-level stacked OLS with event $\times$ ring and event $\times$ calendar-year fixed effects (the ring literature’s convention), clustered by event; and the sale-level specification with hedonic controls. Rings 0–250 m vs 400–1,000 m. The three series coincide within sampling error at every horizon, and the agreement of transaction-weighted and event-weighted estimates indicates the pooled effect is not driven by a few high-volume events.

In the annual robustness suite $Q = 4$ and $H = 3$; the baseline is also estimated at quarterly ($Q = H = 8$) and half-yearly resolutions and with the Callaway–Sant’Anna (2021) group-time estimator $ATT(g, t) = \mathbb{E}[\bar{y}_t - \bar{y}_{g-1} \mid G_g = 1] - \mathbb{E}[\bar{y}_t - \bar{y}_{g-1} \mid \text{never treated}]$, aggregated to event

time; results are similar across estimators and resolutions, and identical with and without tract fixed effects.

4 The eight figures

All figures plot LP-DiD event-study coefficients $\{\beta_{-Q}, \dots, \beta_H\}$ from equations (3)–(4) at annual resolution (pre-window 4, post-window 3), with 95% confidence intervals, cells clustered. Period -1 is the omitted base; the dashed vertical line marks the treatment boundary.




Figure 2: **Baseline event study.** Near-ring (0–250 m) versus far-ring (400–1,000 m) effects on log sale price by year relative to the investor purchase. The headline price bump; flat pre-period coefficients are the design’s core credibility check, and the post-period path shows how quickly the premium appears and whether it persists.



Figure 3: **Investor vs. placebo purchases.** The real event study overlaid with placebo events (non-investor individual buyers in the investor price band, same micro-areas, >3 years from the real event; far rings matched at 400–750 m). A flat placebo says nearby prices respond to *the investor*, not to any high-value transaction; because unmatched investors contaminate the placebo toward finding an effect, a null placebo is conservative evidence of investor specificity.




Figure 4: **Spatial gradient.** Near-ring (0–250 m) and gap-ring (250–400 m) series, each against the same far control. A genuine hyper-local spillover should decay with distance (near > gap > 0); near $\approx$ gap would instead indicate a neighborhood-wide shock partially shared by the control ring.




Figure 5: **Estimated contrast by local event density.** Separate event studies for events with 0–2, 3–25, and >25 other investor events within 1km. Because purchases cluster at the enclave scale (condo buildings, resort communities), other events in dense areas fall disproportionately *inside* the near ring, so the near–far contrast there reflects cumulative local investor exposure rather than a single purchase. The estimated effect is increasing in density and near zero for isolated events: single-purchase channels (comps entry, one renovation, signaling) appear weak on their own, and the price response is a phenomenon of investor *concentration*. Pre-period coefficients are flat in all three density groups, inconsistent with investor waves selecting into already-appreciating pockets.




Figure 6: **Censoring robustness.** Baseline overlaid with events from 2018 onward, for which little time exists for post-period sales to be overwritten by later resales. Agreement with the full sample indicates that CRIM's last-sale-only censoring is not manufacturing the bump.




Figure 7: **Composition checks.** The identical design with property characteristics as outcomes: log assessed structure value (top left), log lot size (top right), condo-unit share (bottom left), vacant-land share (bottom right). Flat lines imply the price bump is appreciation on a stable transaction mix; movements would reveal the stock-upgrading margin of gentrification, changing the interpretation of Figure 2.




Figure 8: **Who transacts: buyers vs. sellers.** Effects on the share of nearby sales with non-Hispanic-named buyers and, separately, sellers (Census 2010 surname classification; corporate and ambiguous names excluded). Buyer share rising faster than seller share means the area is net-importing mainland owners after an investor arrives — the composition shift price indices cannot see.




Figure 9: **Transition types.** Effects on the shares of transactions that are Hispanic-seller→non-Hispanic-buyer versus non-Hispanic-seller→non-Hispanic-buyer (among sales where both parties classify). Rising islander→mainlander transitions indicate conversion of locally held housing to incomers (the displacement margin); rising mainlander→mainlander indicates churn within an already-converted segment — together distinguishing “investors lead conversion” from “investors cluster where conversion already occurred.”

5 The tract-level design and reconciliation (Design 2)

5.1 Data and specification

The tract-level design asks what the ring design structurally cannot: whether investor purchases move *whole neighborhoods*, including the component common to treated areas that near-far contrasts difference out. The outcome panel is the Red Atlas tract $\times$ month transaction file (918 of 981 2020-vintage tracts, 2008–2025), aggregated to tract $\times$ year with the annual price defined as the transaction-count-weighted mean of monthly tract means (composition-noise reduction; the simple mean is a robustness variant). Treatment timing comes from the dated decree-era events: a tract is treated from the year of its first identified investor purchase (235 treated tracts; 1,200 events), with never-treated tracts as controls. Estimation is annual LP-DiD (pre-window 4, post-window 3) on $100 \times \ln(\text{price})$, with and without county fixed effects.

5.2 Results

Prices in treated tracts rise by a pooled +7.6 log points (s.e. 1.8) after the first investor purchase, with flat pre-period coefficients; adding county fixed effects leaves the estimate unchanged (+7.5), and the simple-mean construction gives +6.1 (Figure 10). Splitting treated tracts by investor concentration reproduces the dose pattern at tract scale: tracts with five or more identified purchases (44 tracts) show +18.3 (2.9), tracts with one to four show +3.8 (2.1) (Figure 11). Transaction *volume* falls in treated tracts, but its pre-period coefficients are significantly negative, so parallel trends fails for that outcome and we treat it as descriptive only.




Figure 10: **Tract-level event study: log house price.** Annual LP-DiD on the Red Atlas panel, 235 treated tracts vs never-treated controls; base and county-FE series are indistinguishable and pre-trends are flat.




Figure 11: **Tract-level event study by investor concentration.** Tracts with 5+ identified purchases versus tracts with 1–4, each against never-treated controls: the concentration result at tract scale.

5.3 Reconciliation: the whole vs the sum of the parts

Because the ring design delivers a distance gradient, it implies a *prediction* for each treated tract: the stock-value-weighted average of $\beta(d)$ over the tract’s parcels — the tract-level effect if micro-spillovers were the whole story. Comparing observed tract-DiD estimates to these predictions decomposes the aggregate (Figure 12): tracts with 1–4 purchases sit at or below their prediction (observed +3.8 vs predicted +5.4) — micro-spillovers fully account for them — while the 44 concentration tracts exceed their prediction by roughly 11 log points (observed +18.3 vs predicted +7.5). A market-level channel beyond proximity operates only where investors concentrate.

5.4 Aggregation

Integrating the estimated gradient over the exposure distribution of the island’s improved housing stock (60.0% of parcels, 63.1% of assessed value, lie within 2.5 km of an identified event) yields an island-wide average price effect of +2.1 to +3.0%, or \$6.3–8.0B of housing value at band-specific market/assessed ratios — roughly 6–8% of Puerto Rico’s post-2019 boom (FHFA purchase-only index: +64% 2019–2024). Adding the concentration excess in the 44 high-dose tracts ($\approx$ \$4.3B) raises the identified program footprint to $\approx$ \$11–12B, about 10–11% of the boom. At municipio level, using municipio repeat-sales indices as the local boom denominators, footprints in the top investor municipios are much larger:




Figure 12: **Observed tract effects vs ring-implied predictions.** Observed: LP-DiD pooled post estimates by dose group. Predictions: stock-value-weighted distance-gradient effects over each treated tract's parcels. The all-treated and 1–4-purchase groups sit on their predictions; the 5+ group's excess (≈ 11 log points) is the concentration channel.

Municipio	Local boom '19–'24			Share of boom		
	Growth	\$	Uplift	Synthesis	Ring only	Tract DiD only
Rincón	+148%	\$0.6B	\$0.27B	44.9%	13.3%	54.8%
Dorado	+177%	\$2.5B	\$0.82B	32.2%	11.0%	36.6%
Río Grande	+115%	\$1.8B	\$0.47B	25.6%	8.5%	25.7%
San Juan	+89%	\$16.7B	\$4.13B	24.6%	15.1%	18.9%
Carolina	+91%	\$7.3B	\$1.28B	17.6%	10.4%	12.8%
Guaynabo	+124%	\$5.5B	\$0.85B	15.5%	13.2%	8.5%
Humacao	+121%	\$2.7B	\$0.38B	14.2%	7.4%	12.9%

The dollar boom is the municipio's market stock value times $1 - 1/(1 + g)$, the portion of current value attributable to repeat-sales growth g since 2019. The synthesis column applies the ring gradient to every parcel and adds the concentration excess in 5+-purchase tracts; the design-by-design columns use each design's treatment effects alone — ring only: the distance gradient over all parcels; tract DiD only: the dose-specific tract effects applied to treated tracts' market stock (zero outside treated tracts). The two bracket the synthesis in the hotspots: the ring design misses the tract-wide repricing under concentration, while the tract DiD assigns the whole tract effect to tract stock but nothing to spillovers beyond treated-tract borders.

All aggregates are lower bounds: exposure counts only the 1,200 name-identified purchases, excluding LLC and unmatched investors; the capitalization step applies transaction-price effects to

the full stock; and event effects predating 2019 enter the numerator while the boom clock starts in 2019.

5.5 Mortgage originations: a convincing null

HMDA loan-level data for Puerto Rico (2018–2024; 105,559 originations, aggregated to 2020-vintage tract $\times$ year) show no detectable response of mortgage lending in treated tracts: Poisson (PPML) estimates with tract and year fixed effects are statistically zero for owner-occupied purchases (-0.5% , s.e. 5.4), non-owner-occupied purchases ($+3.0\%$, 11.3), rate/term and cash-out refinancings, and total originations, and are unchanged with county $\times$ year fixed effects. Consistent with decree investors transacting predominantly in cash, the program’s effects appear in prices and buyer composition rather than local credit expansion.

6 Standing caveats

1. **Last-sale censoring** (CRIM): pre-period sales are survivors that never resold; addressed by design and Figure 5.
2. **Site selection**: investors choose appreciating micro-areas; estimates are local spillovers *conditional on* location choice, not “the effect of the Act.”
3. **Match selection**: identified investors skew toward single-property, personal-name buyers; entity (LLC/trust) purchases are missing.
4. **Name proxy**: understates mainland presence; corporate and ambiguous names are excluded from shares.
5. **Assessed values** are on CRIM’s 1957 basis — valid as cross-sectional controls and composition outcomes, never as market values.